

▼ Credit Card Fraud Detection

About This Project

In this project, we aimed to detect fraudulent credit card transactions using a logistic regression model. We began by importing necessary libraries, loading the dataset, and performing initial data exploration. The dataset contained highly imbalanced classes, with a majority of normal transactions (Class 0) and a minority of fraudulent transactions (Class 1).

In future iterations of this project, you could consider the following enhancements

1. Data Loading and Exploration:

- The code begins by importing necessary libraries, including `numpy`, `pandas`, and specific modules from `sklearn`.
- The credit card dataset is loaded using `pd.read_csv('/content/creditcard.csv')`.
- The dataset's structure and basic information are displayed using `.head()`, `.tail()`, and `.info()` methods.
- Missing values are checked using `.isnull().sum()` and are found to be absent in the dataset.
- The distribution of the target variable 'Class' is shown using `.value_counts()`.

2. Data Preparation and Separation:

- The dataset is divided into two subsets: legitimate transactions (`legit`) and fraudulent transactions (`fraud`) based on the 'Class' column.
- Descriptive statistics of the 'Amount' feature are displayed for both legitimate and fraudulent transactions.

3. Feature Engineering and Exploration:

- Summary statistics are shown for the features grouped by the 'Class' variable.

4. Under-Sampling and Balancing:

- Under-sampling is performed to balance the dataset by randomly sampling an equal number of legitimate and fraudulent transactions.
- The concatenated dataset of balanced samples is stored in `new_dataset`.

5. Feature and Target Separation:

- The features (`x`) are extracted by dropping the 'Class' column from the `new_dataset`.
- The target variable (`y`) is extracted, containing the 'Class' column from `new_dataset`.

6. Data Splitting:

- The dataset is split into training and test sets using the `train_test_split` function.

- The dataset is split into training and test sets using the `train_test_split` function.
- The training set is assigned to `x_train` and `y_train`, and the test set is assigned to `x_test` and `y_test`.

7. Model Training and Evaluation:

- A logistic regression model is created using `LogisticRegression()`.
- The model is trained using the training data with `model.fit(x_train, y_train)`.
- The accuracy of the model is evaluated on both the training and test sets using `accuracy_score`.

8. Results:

- The accuracy of the model on the training data is approximately 92.4%.
- The accuracy of the model on the test data is approximately 91.4%.

In conclusion, the provided code focuses on building a simple logistic regression model for credit card fraud detection. The model is trained and evaluated using accuracy as the evaluation metric. While accuracy provides an overview of the model's performance, it is important to consider other evaluation metrics, especially for imbalanced datasets like this one. Further improvements could involve exploring alternative models, utilizing more sophisticated evaluation metrics, and considering other techniques to handle class imbalance.

▼ *Importing Necessary Libraries*

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```

▼ Loading Dataset

```
Credit_Card_data = pd.read_csv('/content/creditcard.csv')
```

```
# first 5 rows
Credit_Card_data.head()
```

1 to 5 of 5 entries Filter ?

index	Time	V1	V2	V3	V4	V5	V6	V7	V8
0	0.0	-1.359807134	-0.072781173	2.536346738	1.378155224	-0.33832077	0.462387778	0.239598554	0.09869
1	0.0	1.191857111	0.266150712	0.166480113	0.448154078	0.060017649	-0.082360809	-0.078802983	0.08510
2	1.0	-1.358354062	-1.340163075	1.773209343	0.379779593	-0.503198133	1.800499381	0.791460956	0.24767
3	1.0	-0.966271712	-0.185226008	1.79299334	-0.863291275	-0.01030888	1.247203168	0.23760894	0.37743
4	2.0	-1.158233093	0.877736755	1.548717847	0.403033934	-0.407193377	0.095921462	0.592940745	-0.27053

```
# Last rows
Credit_Card_data.tail()
```

1 to 5 of 5 entries Filter ?

index	Time	V1	V2	V3	V4	V5	V6	V7	V8
284802	172786.0	-11.88111789	10.07178497	-9.834783457	-2.066655685	-5.364472781	-2.606837331	-4.918215431	7.
284803	172787.0	-0.732788671	-0.05508049	2.035029745	-0.738588584	0.868229399	1.058415272	0.024329696	0.
284804	172788.0	1.91956501	-0.301253846	-3.249639814	-0.557828125	2.63051512	3.031260098	-0.296826527	0.
284805	172788.0	-0.24044005	0.530482513	0.70251023	0.689799168	-0.377961134	0.623707722	-0.686179986	(
284806	172792.0	-0.596412522	-0.189733337	0.703337367	-0.50627124	-0.012545679	-0.649616686	1.577006254	-0.

Show 25 per page



Like what you see? Visit the [data table notebook](#) to learn more about interactive tables.

Warning: Total number of columns (31) exceeds max_columns (20) limiting to first (20) columns.

```
# info
Credit_Card_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
#   Column  Non-Null Count  Dtype
---  -
0    Time      284807 non-null   float64
1    V1         284807 non-null   float64
2    V2         284807 non-null   float64
3    V3         284807 non-null   float64
4    V4         284807 non-null   float64
5    V5         284807 non-null   float64
6    V6         284807 non-null   float64
7    V7         284807 non-null   float64
8    V8         284807 non-null   float64
9    V9         284807 non-null   float64
10   V10        284807 non-null   float64
11   V11        284807 non-null   float64
12   V12        284807 non-null   float64
13   V13        284807 non-null   float64
14   V14        284807 non-null   float64
15   V15        284807 non-null   float64
```

```

16 V16      284807 non-null float64
17 V17      284807 non-null float64
18 V18      284807 non-null float64
19 V19      284807 non-null float64
20 V20      284807 non-null float64
21 V21      284807 non-null float64
22 V22      284807 non-null float64
23 V23      284807 non-null float64
24 V24      284807 non-null float64
25 V25      284807 non-null float64
26 V26      284807 non-null float64
27 V27      284807 non-null float64
28 V28      284807 non-null float64
29 Amount    284807 non-null float64
30 Class     284807 non-null int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB

```

```
# Check missing values
```

```
Credit_Card_data.isnull().sum()
```

```

Time      0
V1        0
V2        0
V3        0
V4        0
V5        0
V6        0
V7        0
V8        0
V9        0
V10       0
V11       0
V12       0
V13       0
V14       0
V15       0
V16       0
V17       0
V18       0
V19       0
V20       0
V21       0
V22       0
V23       0
V24       0
V25       0
V26       0
V27       0
V28       0
Amount    0
Class     0
dtype: int64

```

```
Credit_Card_data['Class'].value_counts()
```

```
0    284315
1       492
Name: Class, dtype: int64
```

▼ This Dataset is Highly Unblanced

▼ 0 = Noramal Transaction

1 = Fraudulent Transaction

```
# Separating the data for analysis
```

```
legit = Credit_Card_data[Credit_Card_data.Class == 0]
fraud = Credit_Card_data[Credit_Card_data.Class == 1]
```

```
print(legit.shape)
print(fraud.shape)
```

```
(284315, 31)
(492, 31)
```

```
# Statical Data
```

```
legit.Amount.describe()
```

```
count    284315.000000
mean       88.291022
std       250.105092
min         0.000000
25%         5.650000
50%        22.000000
75%        77.050000
max      25691.160000
Name: Amount, dtype: float64
```

```
fraud.Amount.describe()
```

```
count      492.000000
mean      122.211321
std       256.683288
min         0.000000
25%         1.000000
50%         9.250000
```

```
75%      105.890000
max      2125.870000
Name: Amount, dtype: float64
```

```
# compare the values for the both
Credit_Card_data.groupby('Class').mean()
```

1 to 2 of 2 entries Filter ?

Class	Time	V1	V2	V3	V4
0	94838.20225805884	0.008257515901953819	-0.0062708574165133745	0.012170917030775724	-0.007859867819657
1	80746.80691056911	-4.771948441388211	3.6237781019126016	-7.033281048536585	4.542029104436

Show 25 per page



Like what you see? Visit the [data table notebook](#) to learn more about interactive tables.

Warning: Total number of columns (30) exceeds max_columns (20) limiting to first (20) columns.

Under-Sampling

```
#
legit_sample = legit.sample(n=492)
```

```
# Concatenating two Data
```

```
new_dataset = pd.concat([legit_sample,fraud],axis=0)
```

```
new_dataset.head()
```

1 to 5 of 5 entries Filter ?

index	Time	V1	V2	V3	V4	V5	V6	V7
249671	154524.0	0.138621231	0.504223165	0.226955681	-0.628245845	0.807858252	-1.095610497	1.290402334
193201	130035.0	2.135805959	-1.01803093	-0.757191627	-0.766406588	-1.094454941	-0.920283835	-0.811455336
131528	79634.0	1.211782313	0.041224299	0.329701555	1.144392329	-0.20881684	-0.072755671	-0.087446824
97391	66199.0	-1.077886975	1.098046354	1.18535212	-0.388477331	0.871209159	0.052171301	0.929343359
118899	75249.0	1.081609486	-1.797639018	0.569293681	-1.134711787	-2.099889495	-0.743406984	-0.998303228

Show 25 per page



Like what you see? Visit the [data table notebook](#) to learn more about interactive tables.

Warning: Total number of columns (31) exceeds max_columns (20) limiting to first (20) columns.

```
new_dataset.tail()
```

1 to 5 of 5 entries Filter ?

index	Time	V1	V2	V3	V4	V5	V6	V7	
279863	169142.0	-1.927883321	1.125652664	-4.518330641	1.749292533	-1.566487292	-2.010494231	-0.882849831	0.6
280143	169347.0	1.378558997	1.289380937	-5.004246784	1.411849842	0.442580636	-1.326535934	-1.413169956	0.2
280149	169351.0	-0.676142671	1.126366062	-2.213699523	0.468308388	-1.120541044	-0.003346296	-2.234739296	1.2
281144	169966.0	-3.113831607	0.585864172	-5.399730211	1.817092473	-0.840618466	-2.943547791	-2.20800192	1.0
281674	170348.0	1.991976096	0.158475887	-2.583440645	0.408669993	1.151147061	-0.096694744	0.223050267	-0.0

Show per page



Like what you see? Visit the [data table notebook](#) to learn more about interactive tables.

Warning: Total number of columns (31) exceeds max_columns (20) limiting to first (20) columns.

```
new_dataset['Class'].value_counts()
```

```
0    492
1    492
Name: Class, dtype: int64
```

```
new_dataset.groupby('Class').mean()
```

1 to 2 of 2 entries Filter ?

Class	Time	V1	V2	V3	V4
0	94868.6361788618	0.044826651589430895	-0.06486907317479675	-0.04177646449390244	-0.11030721061991869
1	80746.80691056911	-4.771948441388211	3.6237781019126016	-7.033281048536585	4.542029104436992

Show per page



Like what you see? Visit the [data table notebook](#) to learn more about interactive tables.

Warning: Total number of columns (30) exceeds max_columns (20) limiting to first (20) columns.

▸ Splitting the data into features & Target

```
x = new_dataset.drop(columns='Class', axis=1)
y = new_dataset['Class']
```

```
print(x)
```

```
      Time      V1      V2      V3      V4      V5      V6  \
249671  154524.0  0.138621  0.504223  0.226956 -0.628246  0.807858 -1.095610
193201  130035.0  2.135806 -1.018031 -0.757192 -0.766407 -1.094455 -0.920284
```

131528	79634.0	1.211782	0.041224	0.329702	1.144392	-0.208817	-0.072756
97391	66199.0	-1.077887	1.098046	1.185352	-0.388477	0.871209	0.052171
118899	75249.0	1.081609	-1.797639	0.569294	-1.134712	-2.099889	-0.743407
...
279863	169142.0	-1.927883	1.125653	-4.518331	1.749293	-1.566487	-2.010494
280143	169347.0	1.378559	1.289381	-5.004247	1.411850	0.442581	-1.326536
280149	169351.0	-0.676143	1.126366	-2.213700	0.468308	-1.120541	-0.003346
281144	169966.0	-3.113832	0.585864	-5.399730	1.817092	-0.840618	-2.943548
281674	170348.0	1.991976	0.158476	-2.583441	0.408670	1.151147	-0.096695

	V7	V8	V9	...	V20	V21	V22	\
249671	1.290402	-0.503858	0.022366	...	-0.046210	-0.145491	-0.215729	
193201	-0.811455	-0.148290	0.122525	...	-0.074043	0.202852	0.567046	
131528	-0.087447	0.010871	0.655480	...	-0.127124	-0.319999	-0.820429	
97391	0.929343	-0.004249	-0.564004	...	0.079059	-0.514485	-1.465811	
118899	-0.998303	-0.171101	-1.548716	...	0.089188	0.064507	0.088093	
...	
279863	-0.882850	0.697211	-2.064945	...	1.252967	0.778584	-0.319189	
280143	-1.413170	0.248525	-1.127396	...	0.226138	0.370612	0.028234	
280149	-2.234739	1.210158	-0.652250	...	0.247968	0.751826	0.834108	
281144	-2.208002	1.058733	-1.632333	...	0.306271	0.583276	-0.269209	
281674	0.223050	-0.068384	0.577829	...	-0.017652	-0.164350	-0.295135	

	V23	V24	V25	V26	V27	V28	Amount
249671	0.235163	-0.110294	-1.435347	-0.055238	0.037253	0.045597	26.99
193201	0.168626	-0.088614	-0.190652	-0.170339	-0.011791	-0.051808	32.00
131528	-0.072292	-0.482841	0.577632	-0.493104	0.030474	0.021828	28.27
97391	0.038850	-0.982764	0.157070	-0.006196	-0.002640	0.129799	14.65
118899	-0.185738	0.369997	0.245524	-0.124256	0.008085	0.063739	228.00
...
279863	0.639419	-0.294885	0.537503	0.788395	0.292680	0.147968	390.00
280143	-0.145640	-0.081049	0.521875	0.739467	0.389152	0.186637	0.76
280149	0.190944	0.032070	-0.739695	0.471111	0.385107	0.194361	77.89
281144	-0.456108	-0.183659	-0.328168	0.606116	0.884876	-0.253700	245.00
281674	-0.072173	-0.450261	0.313267	-0.289617	0.002988	-0.015309	42.53

[984 rows x 30 columns]

```
print(y)
```

```

249671    0
193201    0
131528    0
97391     0
118899    0
..
279863    1
280143    1
280149    1
281144    1
281674    1
Name: Class, Length: 984, dtype: int64

```


▼ Split the data into Training data & Test

```
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.2, stratify=y, random_state = 2)
```

```
print(x.shape, x_train.shape, x_test.shape)
```

```
(984, 30) (787, 30) (197, 30)
```

Model Training

▼ Logistic Regression

```
model = LogisticRegression()
```

```
# Train Data  
model.fit(x_train, y_train)
```

```
▼ LogisticRegression  
LogisticRegression()
```

▼ Model Evaluation

Accuracy Score

```
# Accuracy on Training Data  
x_train_prediction = model.predict(x_train)  
training_data_accuracy = accuracy_score(x_train_prediction, y_train)
```

```
print('Accuracy on Training data: ',training_data_accuracy)
```

```
Accuracy on Training data: 0.9237611181702668
```

```
# accuracy on test Data  
x_test_prediction = model.predict(x_test)  
test_data_accuracy = accuracy_score(x_test_prediction, y_test)
```

```
print('Accuracy on Test data: ',test_data_accuracy)
```

```
Accuracy on Test data:  0.9137055837563451
```

✓ 0s completed at 1:29 PM

